

ML ASSIGNMENT 3

KABIR CHATURVEDI , J011

IMPORTING:

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

UNIVARIATE REGRESSION:

```
In [2]: data = pd.read_csv('ex1data1.txt',header=None)
```

```
In [3]: data.head()
```

Out[3]:

	0	1
0	6.1101	17.5920
1	5.5277	9.1302
2	8.5186	13.6620
3	7.0032	11.8540
4	5.8598	6.8233

```
In [4]: # NOW WE NAME THE COLUMNS
data.columns = ['Population','Profit']
```

```
In [5]: data.head()
```

Out[5]:

	Population	Profit
0	6.1101	17.5920
1	5.5277	9.1302
2	8.5186	13.6620
3	7.0032	11.8540
4	5.8598	6.8233

```
In [6]: data.isna().sum()
```

Out[6]:

Population	0
Profit	0
dtype: int64	

```
In [7]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 97 entries, 0 to 96
Data columns (total 2 columns):
 # Column   Non-Null Count  Dtype
---  -
0  Population  97 non-null    float64
1  Profit     97 non-null    float64
dtypes: float64(2)
memory usage: 1.6 KB
```

```
In [8]: data.describe()
```

Out[8]:

	Population	Profit
count	97.000000	97.000000
mean	8.159800	5.839135
std	3.869884	5.510262
min	5.026900	-2.680700
25%	5.707700	1.986900
50%	6.589400	4.562300
75%	8.578100	7.046700
max	22.203000	24.147000

```
In [9]: X = data.iloc[:, :-1].values
y = data.iloc[:, 1].values.reshape(-1,1)
```

```
In [10]: X.shape, y.shape
```

```
Out[10]: ((97, 1), (97, 1))
```

```
In [11]: m = y.size
ones = np.ones((m,1)) # ADDING A MATRIX FULL OF ONES
X1 = np.c_[ones,X]
```

```
In [12]: X1.shape, y.shape
```

```
Out[12]: ((97, 2), (97, 1))
```

```
In [13]: n = X1.shape[1]
theta1 = np.zeros((n,1)) # ADDING A MATRIX FULL OF ZEROES
theta1.shape
```

```
Out[13]: (2, 1)
```

```
In [14]: def calCost(error):
m = error.shape[0]
squared_error = np.dot(error.T,error)
return squared_error[0][0]/m
```


```
In [15]: def GradientDescent(X,y,theta,alpha,epochs):
cost_history = []
theta_history = [theta]
m = y.size
for i in range(epochs):
h_theta = np.dot(X,theta)
error = h_theta - y

cost_history.append(calCost(error))
diff = np.dot(X.T, error)/m
theta = theta - (alpha*diff)
theta_history.append(theta)
return cost_history, theta_history
```

```
In [16]: cost_history, theta_history = GradientDescent(X=X1,y=y,theta=theta1,alpha=0.001,epochs=2000)
```

- Plotting the cost history:

```
In [17]: plt.figure(figsize=(8,4),dpi=200)
plt.plot(cost_history);
```



```
In [18]: theta_history[-1] # the latest value of theta
```

```
Out[18]: array([[ -1.12369018],
[ 0.91454787]])
```

```
In [19]: from sklearn.linear_model import LinearRegression
```

```
In [20]: model = LinearRegression()
```

```
In [21]: model.fit(X1,y)
```

```
Out[21]: LinearRegression()
```

```
In [22]: model.intercept_ , model.coef_
```

```
Out[22]: (array([-3.89578088]), array([[0.          , 1.19383364]]))
```

TIME FOR MULTIVARIATE REGRESSION:

```
In [23]: new_data = pd.read_csv('ex1data2.txt',header=None)
```

```
In [24]: new_data.head()
```

Out[24]:

	0	1	2
0	2104	3	399900
1	1600	3	329900
2	2400	3	369000
3	1416	2	232000
4	3000	4	539900

```
In [25]: new_data.isna().sum()
```

Out[25]:

0	0
1	0
2	0
dtype: int64	

```
In [26]: new_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 47 entries, 0 to 46
Data columns (total 3 columns):
 # Column   Non-Null Count  Dtype
---  -
0  0         47 non-null    int64
1  1         47 non-null    int64
2  2         47 non-null    int64
dtypes: int64(3)
memory usage: 1.2 KB
```

```
In [27]: new_data.describe()
```

Out[27]:

	0	1	2
count	47.000000	47.000000	47.000000
mean	2000.680851	3.170213	340412.659574
std	794.702354	0.760982	125039.899586
min	852.000000	1.000000	169900.000000
25%	1432.000000	3.000000	249900.000000
50%	1888.000000	3.000000	299900.000000
75%	2269.000000	4.000000	384450.000000
max	4478.000000	5.000000	699900.000000

```
In [28]: new_data.columns = ['Size of the house(Sq. Feet)', 'No. of Bedrooms', 'Price']
```

```
In [29]: new_data.head()
```

Out[29]:

	Size of the house(Sq. Feet)	No. of Bedrooms	Price
0	2104	3	399900
1	1600	3	329900
2	2400	3	369000
3	1416	2	232000
4	3000	4	539900

Creating a function to normalize the data:

```
In [30]: def normalize(dataframe):
df=dataframe.copy()
for col in df.columns:
df[col] = (df[col]-df[col].mean())/df[col].std()
```

```
In [31]: new_data.columns
```

```
Out[31]: Index(['Size of the house(Sq. Feet)', 'No. of Bedrooms', 'Price', dtype='object')
```

```
In [32]: new_data = normalize(new_data)
```

```
In [33]: X2 = new_data.iloc[:, :-1].values
y2 = new_data.iloc[:, 1].values.reshape(-1,1)
```

```
In [34]: X2.shape, y2.shape
```

```
Out[34]: ((47, 2), (47, 1))
```

```
In [35]: m = y2.size
ones = np.ones((m,1)) # ADDING A MATRIX FULL OF ONES
X2 = np.c_[ones,X2]
```

```
In [36]: X2.shape, y2.shape
```

```
Out[36]: ((47, 3), (47, 1))
```


```
In [37]: n = X2.shape[1]
theta2 = np.zeros((n,1)) # ADDING A MATRIX FULL OF ZEROES
theta2.shape
```

```
Out[37]: (3, 1)
```

```
In [38]: new_cost_history, new_theta_history = GradientDescent(X2,y2,theta2,alpha=0.001,epochs=2000)
```

- Plotting the cost history for the SECOND model:

```
In [40]: plt.figure(figsize=(8,4),dpi=200)
plt.plot(new_cost_history);
```



```
In [41]: new_theta_history[-1] #latest theta
```

```
Out[41]: array([[ -1.19667869e-16],
[ 6.67841350e-01],
[ 1.25393822e-01]])
```

```
In [42]: new_model = LinearRegression()
```

```
In [43]: new_model.fit(X2,y2)
```

```
Out[43]: LinearRegression()
```

```
In [44]: new_model.intercept_ , new_model.coef_
```

```
Out[44]: (array([-1.15685754e-16]), array([[ 0.          , 0.88476599, -0.05317882]]))
```